

Tennis/Pickleball Single-Camera Tracking System

Progress Report & Task Allocation

Course: Computer Vision

Instructor: Nguyen Duc Dung

Team Members:

1. Pham Hoang An 2. Dao Duy Anh 3. Duong Vu Duc
4. Nguyen Duc Dat 5. Bui Dinh Nguyen Minh

GitHub: <https://github.com/AnHgPham/cv>

Report date: February 23, 2026

Contents

1 Abstract	3
2 Repository and Project Structure	3
2.1 Directory Layout	3
3 System Architecture	4
3.1 Pipeline Flow Diagram	4
3.2 Per-Frame Processing Steps	4
4 Module Details	5
4.1 Module 1: Court Detection & Homography	5
4.1.1 Technical Approach	5
4.1.2 Additional Components	5
4.1.3 Known Constants	5
4.1.4 Applied Knowledge	5
4.2 Module 2: Object Detection	5
4.2.1 Detectors Implemented	5
4.2.2 Detection Priority	6
4.2.3 Applied Knowledge	6
4.3 Module 3: Object Tracking	6
4.3.1 Ball Tracking	6
4.3.2 Player Tracking	6
4.3.3 Applied Knowledge	7
4.4 Module 4: 3D Trajectory Reconstruction & In/Out Classification	7
4.4.1 Trajectory Reconstruction	7
4.4.2 Bounce Detection	7
4.4.3 In/Out Classification	7
4.4.4 Applied Knowledge	7

4.5	Module 5: Pipeline Integration & Visualization	7
4.5.1	Pipeline Orchestration	7
4.5.2	Visualization Components	8
4.5.3	Data Preprocessing Utilities	8
4.5.4	Applied Knowledge	8
5	Bugs Found and Fixed During Development	8
6	Remaining Issue: Incorrect Homography	9
6.1	Symptom	9
6.2	Root Cause Analysis	9
6.3	Impact	9
7	Proposed Fixes	9
7.1	Fix A: Robust Corner Selection for Homography (Module 1)	9
7.2	Fix B: Show Players on Mini-map (Module 5)	10
7.3	Fix C: Court-Line Overlay on Video Frame (Module 5)	10
8	Task Allocation	10
9	Dependency Graph and Critical Path	13
10	Milestones	13
11	Test Data and Validation	13
11.1	Primary Test Video	13
11.2	Latest Pipeline Results (Before Homography Fix)	14
11.3	Validation Plan	14
12	Environment and Reproduction	14
12.1	Environment	14
12.2	Installation	14
12.3	Running the Pipeline	14
12.4	Inspecting a Single Frame	15
13	Conclusion	15

1 Abstract

This project implements a complete single-camera computer vision pipeline for tennis and pickleball video analysis. The system processes monocular video (25–30 fps) and produces:

- Real-time ball trajectory detection and tracking,
- Player detection with persistent identity assignment,
- Court-line detection and image-to-court homography estimation,
- Physics-based 3D trajectory reconstruction of the ball,
- Bounce detection with in/out classification,
- Annotated output video with bird’s-eye mini-map and heatmaps.

This report presents the current implementation status, documents bugs identified and fixed during development, highlights remaining issues (primarily an incorrect homography due to multi-court interference), and defines a task allocation for a 5-member team to deliver a stable, evaluated system.

2 Repository and Project Structure

The full source code is hosted on GitHub:

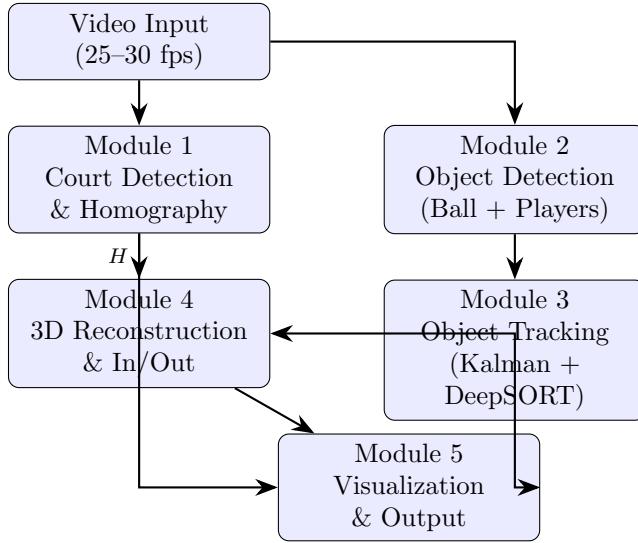
- Repository: <https://github.com/AnHgPham/cv>
- Project root: cv/tennis-pickleball-tracker/

2.1 Directory Layout

```
tennis-pickleball-tracker/
|-- src/
|   |-- court_detection.py      # Module 1 (~25 KB, ~720 lines)
|   |-- object_detection.py    # Module 2 (~35 KB, ~400 lines)
|   |-- object_tracking.py     # Module 3 (~24 KB, ~540 lines)
|   |-- trajectory_3d.py       # Module 4a (~31 KB, ~680 lines)
|   |-- in_out_classifier.py   # Module 4b (~9 KB, ~250 lines)
|   |-- pipeline.py            # Module 5a (~36 KB, ~940 lines)
|   |-- visualization.py       # Module 5b (~21 KB, ~630 lines)
|   '-- __init__.py
|-- notebooks/
|   |-- 01_data_exploration.ipynb
|   |-- 02_court_detection.ipynb
|   |-- 03_ball_detection.ipynb
|   |-- 04_tracking.ipynb
|   '-- 05_3d_reconstruction.ipynb
|-- configs/court_config.yaml
|-- data/raw/                  # Input videos (not in git)
|-- outputs/                   # Generated results (not in git)
|-- run_pipeline.py            # CLI entry point
|-- test_fix.py                # Single-frame debug script
`-- report.tex                 # This report
```

3 System Architecture

3.1 Pipeline Flow Diagram



3.2 Per-Frame Processing Steps

For every frame t in the input video:

1. **Court Detection (Module 1):** Detect white court lines via HSV thresholding, extract edges with Canny, find line segments with Hough transform, compute line intersections, select 4 court corners, and estimate a 3×3 homography matrix H that maps pixel coordinates (u, v) to real-world court coordinates (x_c, y_c) in metres.
2. **Object Detection (Module 2):** Run YOLOv8 on the frame to detect persons and sports balls. Supplement with a classical contour-based ball detector for small tennis balls that YOLO may miss. Apply non-maximum suppression (NMS) to remove duplicate detections.
3. **Player Filtering (Module 5):** Score each player detection by (a) x-distance from frame centre (weight 0.60), (b) perspective-consistent size (weight 0.25), (c) YOLO confidence (weight 0.15). Hard-reject detections whose centre exceeds 40% of frame width from the frame centre. Keep top k (default $k = 4$ for doubles).
4. **Tracking (Module 3):** Update the ball Kalman filter with the new detection (or predict if no detection). Use Lucas-Kanade optical flow as a secondary motion estimate. Update the player tracker (DeepSORT or IoU fallback) to maintain persistent IDs.
5. **3D Reconstruction (Module 4):** Project the tracked ball position through H to obtain court coordinates. Estimate ball height using a physics model (gravity $g = 9.81 \text{ m/s}^2$, air drag). Detect bounces from trajectory direction changes. Classify each bounce as in or out by checking whether (x_c, y_c) falls within the standard court rectangle.
6. **Visualization (Module 5):** Draw bounding boxes, ball trajectory trail, info overlay, and compose the output frame with a side-panel mini-map and frame metadata.

4 Module Details

4.1 Module 1: Court Detection & Homography

File: `src/court_detection.py`

Owner: Pham Hoang An

4.1.1 Technical Approach

1. **Pre-processing:** Convert BGR → HSV; threshold for white pixels (H: 0–180, S: 0–50, V: 180–255); morphological dilation (2 iterations) and erosion (1 iteration) with a 5×5 rectangular kernel.
2. **Edge Detection:** Canny with thresholds (50, 150).
3. **Line Detection:** Probabilistic Hough Transform ($\rho = 1$, $\theta = \pi/180$, threshold= 100, minLineLength= 50, maxLineGap= 30).
4. **Line Classification:** Separate into horizontal and vertical groups based on angle ($\pm 30^\circ$ from horizontal/ vertical). Merge similar lines within 30 px distance.
5. **Intersection Computation:** Compute all pairwise intersections of horizontal and vertical lines.
6. **Corner Selection:** Compute centroid of all intersection points; assign each point to a quadrant (TL, TR, BL, BR) relative to the centroid; pick one representative per quadrant.
7. **Homography:** Compute H via `cv2.findHomography` with RANSAC (reproj. threshold = 5.0 px) mapping 4 image corners → standard court corners (tennis: 23.77×10.97 m).

4.1.2 Additional Components

- **SIFT Court Matcher:** Computes frame-to-frame homography using SIFT keypoints + FLANN matcher + Lowe's ratio test, enabling camera-motion compensation.
- **Deep Court Detector (scaffolded):** CNN-based keypoint regression using a ResNet-18 backbone. Not yet trained.

4.1.3 Known Constants

Standard tennis court keypoints are defined in `TENNIS_COURT_KEYPOINTS` (21 points in metres) and `TENNIS_COURT_CORNERS` (4 corners). Pickleball court keypoints are also defined.

4.1.4 Applied Knowledge

Image processing (colour spaces, morphology), Canny edge detection, Hough transform, SIFT feature matching, homography estimation (DLT + RANSAC), convolutional neural networks.

4.2 Module 2: Object Detection

File: `src/object_detection.py`

Owner: Dao Duy Anh

4.2.1 Detectors Implemented

1. **YOLODetector:** Wraps Ultralytics YOLOv8.

- Supports both COCO-pretrained models (class 0 = person, class 32 = sports ball) and custom-trained models (user-defined class map).
 - Configurable confidence threshold (default 0.25), device (CPU/CUDA).
 - Detected 9 persons in a typical frame of our test video.
2. **ClassicalBallDetector:** A multi-stage pipeline for detecting small tennis balls that YOLO often misses:
 - HSV colour filtering for yellow/green tennis balls (H: 20–50, S: 50–255, V: 120–255).
 - Contour extraction with circularity filter ($4\pi A/P^2 > 0.6$).
 - Size filtering (radius 3–30 px).
 - Motion detection via frame differencing.
 - Court-region filtering (optional, uses court polygon).
 3. **TrackNetDetector (scaffolded):** A specialised heatmap regression network that uses 3 consecutive frames as input to detect fast-moving balls.
 4. **ClassicalPlayerDetector:** HOG + SVM or Haar cascade baseline for player detection.
 5. **FasterRCNNDetector (scaffolded):** Two-stage detector using torchvision’s pre-trained Faster R-CNN.

4.2.2 Detection Priority

The pipeline tries detectors in this order: classical → YOLO → TrackNet. The first detector to return a ball detection “wins” for that frame.

4.2.3 Applied Knowledge

Feature engineering (HOG, Haar), deep object detection (YOLO, Faster R-CNN), heatmap regression (TrackNet), non-maximum suppression, colour space analysis, contour geometry.

4.3 Module 3: Object Tracking

File: `src/object_tracking.py`

Owner: Duong Vu Duc

4.3.1 Ball Tracking

- **Kalman Filter:** State vector $\mathbf{x} = [x, y, v_x, v_y]^T$. Constant-velocity motion model. Configurable process noise (default 5.0) and measurement noise (default 2.0).
- **Optical Flow:** Lucas-Kanade sparse optical flow around the last known ball position as a secondary motion estimate.
- **Fusion:** Weighted combination of Kalman prediction and optical flow result. If no detection for > 10 frames, the track is dropped.

4.3.2 Player Tracking

- **DeepSORT (preferred):** Uses deep appearance features + Kalman prediction + Hungarian algorithm for data association. Requires `deep-sort-realtim` package.
- **IoU Tracker (fallback):** Simple IoU-based matching between consecutive frames. Currently active because `deep-sort-realtim` is not installed. Track IDs may be less stable under occlusion or missed detections.

4.3.3 Applied Knowledge

Kalman filtering (linear state estimation), Lucas-Kanade optical flow, Hungarian algorithm (optimal assignment), deep metric learning (DeepSORT appearance model), IoU-based data association.

4.4 Module 4: 3D Trajectory Reconstruction & In/Out Classification

Files: `src/trajectory_3d.py`, `src/in_out_classifier.py`

Owner: Nguyen Duc Dat

4.4.1 Trajectory Reconstruction

- Projects 2D ball pixel positions (u, v) to court coordinates (x_c, y_c) using the homography H .
- Estimates ball height $z(t)$ using a physics model: $z(t) = z_0 + v_{z0}t - \frac{1}{2}gt^2$ with optional air-drag correction.
- Uses an Extended Kalman Filter for 3D state smoothing.

4.4.2 Bounce Detection

- Detects frames where the vertical velocity v_z changes sign (trajectory direction reversal).
- Filters spurious bounces using minimum height and time constraints.

4.4.3 In/Out Classification

- Checks whether the bounce point (x_c, y_c) falls within the standard court rectangle (with line-width tolerance).
- Outputs a confidence score based on distance to the nearest boundary.
- ML-enhanced mode (scaffolded): uses nearby trajectory features for more robust classification.

4.4.4 Applied Knowledge

Projective geometry, coordinate transformations, physics-based modelling (projectile motion, air drag), Extended Kalman filter, geometric classification.

4.5 Module 5: Pipeline Integration & Visualization

Files: `src/pipeline.py`, `src/visualization.py`

Owner: Bui Dinh Nguyen Minh

4.5.1 Pipeline Orchestration

- `TennisPickleballPipeline` class manages all module instances and state (homography, court polygon, ball history).
- `PipelineConfig` supports YAML config files and CLI arguments for all parameters.
- Court detection is re-run every 30 frames to handle gradual camera drift.
- Player filtering uses a multi-factor scoring system:

$$S = 0.15 \cdot c_{\text{conf}} + 0.25 \cdot c_{\text{size}} + 0.60 \cdot c_{\text{court}}$$

where c_{court} penalises players far from the frame centre (proxy for main-court membership).

4.5.2 Visualization Components

- **FrameAnnotator:** Draws bounding boxes (colour-coded by player ID), ball trajectory trail with fade effect, bounce markers (green circle = in, red X = out), and info overlay.
- **MiniMap:** Renders a 300×600 px bird's-eye view of the court with proper scaling (23.77×10.97 m for tennis). Supports drawing ball position, player positions, ball trail, and bounce markers. Currently only receives ball position from the pipeline.
- **HeatmapGenerator:** Accumulates ball and player positions in a discretised grid (0.1 m resolution), applies Gaussian smoothing ($\sigma = 2$), and renders via matplotlib with court line overlay.
- **CompositeFrameBuilder:** Combines the main annotated frame (1280×720) with a side panel containing the mini-map (200×400) and text info.

4.5.3 Data Preprocessing Utilities

- Frame extraction from video at configurable intervals.
- Resize and normalise frames for TrackNet / YOLO input.
- Dataset splitting (train/val/test) with configurable ratios.
- Annotation format conversion (to YOLO format).

4.5.4 Applied Knowledge

Software engineering (modular pipeline design, configuration management), video I/O (OpenCV VideoCapture/VideoWriter), data visualisation, image composition.

5 Bugs Found and Fixed During Development

During the development and testing phase, the following bugs were identified and resolved:

1. **YOLO Class Mapping Error (Module 2).** The original code assumed class indices {0: “ball”, 1: “player”}, which is correct for custom-trained models but *wrong* for COCO-pretrained YOLOv8, where class 0 = “person” and class 32 = “sports ball”.

Fix: Added an `is_custom_model` parameter to `YOLODetector`. In COCO mode, the detector uses {0: ‘‘player’’, 32: ‘‘ball’’} and filters for classes [0, 32].

2. **Player Detection Included Adjacent Courts (Module 5).** With 9 YOLO person detections in a typical frame, the system initially displayed all of them, including spectators and players on adjacent courts.

Fix: Implemented a multi-factor scoring system (Section 3) that scores each detection by proximity to the frame centre, perspective-consistent size, and confidence. A hard cutoff at 40% of frame width rejects far-edge detections. Top $k = 4$ are kept for doubles.

3. **Ball Detection Missed Small Tennis Balls (Module 2).** YOLOv8 with COCO weights frequently missed the tennis ball (small, fast-moving, ~ 10 px diameter).

Fix: Rewrote `ClassicalBallDetector` with contour analysis, circularity filtering, motion detection via frame differencing, and court-region filtering. The pipeline now tries classical detection first, then falls back to YOLO.

4. **Import Error in Tracking Module (Module 3).** A relative import of `_compute_iou` failed when running the module standalone.

Fix: Added try/except fallback for both relative and absolute import paths.

5. **Pipeline Ball Label Mismatch (Module 5).** The pipeline’s ball filter checked for “ball” but COCO models return “sports ball”.

Fix: Updated the filter to accept both: (“ball”, “sports ball”).

6 Remaining Issue: Incorrect Homography

6.1 Symptom

When the input video contains multiple adjacent courts (as in our test video), the mini-map ball position is misplaced and all bounces are classified as OUT, even though the ball visually lands inside the main court.

6.2 Root Cause Analysis

1. **Corner selection picks wrong court.** The classical court detector finds line intersections from *all visible courts* (15 intersection points detected; many at extreme coordinates such as $x = -3562$ or $x = 3313$). The centroid-quadrant heuristic in `_select_court_corners` mixes intersections from the main court with those from adjacent courts, producing a degenerate homography.
2. **Homography validation is absent.** There is no sanity check on the computed homography. Projecting standard court corners through H^{-1} back to image space yields a polygon with vertices at $(x_{\min}, y_{\min}) = (700, -325)$ and $(x_{\max}, y_{\max}) = (2910, 1080)$, far exceeding the 1920×1080 frame. The near-side player’s foot position (583, 1075) projects to court coordinates $(-6.56, 5.19)$ instead of a point inside $[0, 23.77] \times [0, 10.97]$.
3. **Mini-map does not show player positions.** The `MiniMap.render()` method supports `player_court_positions` but the pipeline integration code only passes `ball_court_pos`, so the mini-map displays only a (mis-mapped) ball dot on a static court diagram.

6.3 Impact

- All 4 detected bounces are classified as OUT (incorrect).
- Ball heatmap accumulates at wrong court positions.
- Mini-map is essentially static and uninformative.
- Modules 4 and 5 cannot be properly validated until the homography is corrected.

7 Proposed Fixes

7.1 Fix A: Robust Corner Selection for Homography (Module 1)

1. **Intersection filtering:** Remove all intersections outside the frame ($\pm 5\%$ margin). Cluster remaining points and select the largest cluster as the main court.
2. **Aspect-ratio constraint:** Among candidate 4-point subsets, select the one whose projected rectangle best matches the standard court aspect ratio (≈ 2.17).

3. **Reprojection sanity check:** After computing H , project standard court corners back to image space. Reject if: (a) any projected point is more than 20% outside the frame, (b) the projected polygon is not convex, or (c) reprojection error exceeds 10 px.
4. **Temporal stabilisation:** Cache the last valid H and only replace it when a new H passes all sanity checks with better quality. Optionally average H matrices over a sliding window.

7.2 Fix B: Show Players on Mini-map (Module 5)

1. For each tracked player, compute foot point = bottom-centre of bounding box: $(\frac{x_1+x_2}{2}, y_2)$.
2. Project foot point through H to court coordinates.
3. Pass the list of court positions to `MiniMap.render(player_court_positions=...)`.
4. Guard against bad projections (NaN, out-of-court-range) by clamping or skipping.

7.3 Fix C: Court-Line Overlay on Video Frame (Module 5)

1. Project all 21 standard tennis court keypoints through H^{-1} to image coordinates.
2. Draw the projected lines as a semi-transparent overlay on the video frame, providing immediate visual feedback on homography quality.

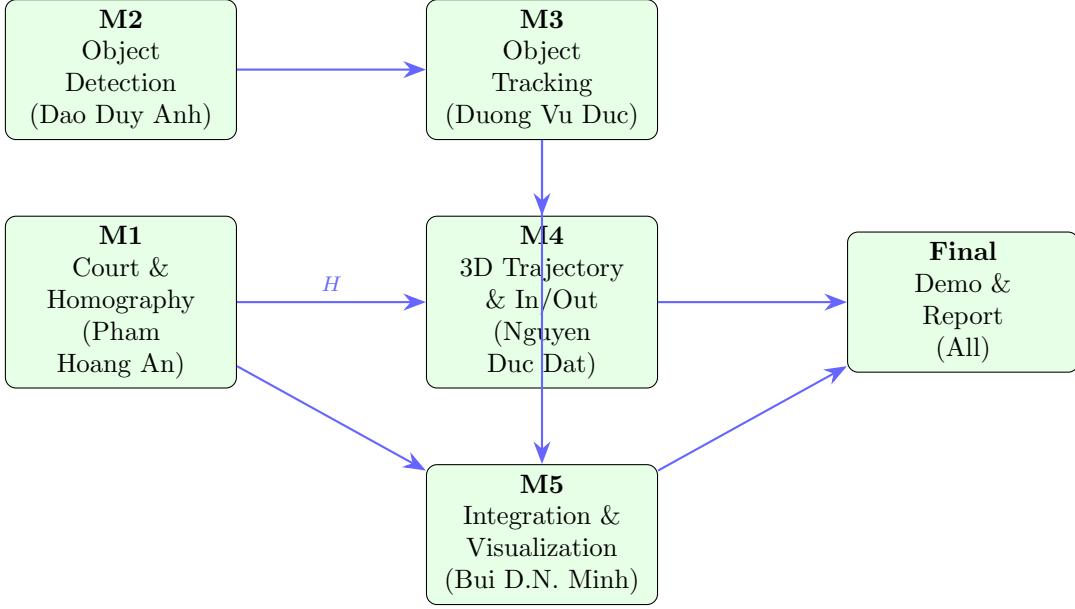
8 Task Allocation

Member	Module	Responsibilities & Deliverables	Status
Pham Hoang An	Module 1: Court Detection & Homography	<ul style="list-style-type: none"> • Filter intersections: remove out-of-frame points; cluster remaining points to isolate the main court. • Implement aspect-ratio constraint for corner selection. • Add reprojection sanity check: reject degenerate homographies. • Temporal stabilisation: cache last valid H, smooth updates. • (Optional) Train the scaffolded ResNet-18 court keypoint detector on labelled data. • Write unit tests for single-court and multi-court scenarios. • Deliverables: updated <code>court_detection.py</code>, before/after comparison images, test report. • Knowledge: Image processing, Hough Transform, RANSAC, homography, SIFT, CNN. 	In progress

Member	Module	Responsibilities & Deliverables	Status
Dao Duy Anh	Module 2: Object Detection	<ul style="list-style-type: none"> Benchmark ball detection: classical vs. YOLOv8 (COCO) vs. TrackNet. Report detection rate, precision, recall on 100+ annotated frames. Tune classical detector parameters (HSV range, contour circularity threshold, size limits). Reduce false-positive ball detections using court polygon constraint (reject candidates outside court area). Evaluate player detection: measure precision/recall, analyse failure cases (occlusion, small far-side players). (Optional) Fine-tune YOLOv8 on tennis-specific dataset or implement TrackNet training. Deliverables: detection evaluation report with metrics, optimised config, updated <code>object_detection.py</code>. Knowledge: HOG, Haar cascades, YOLO, Faster R-CNN, TrackNet, NMS, colour-space filtering. 	In progress
Duong Vu Duc	Module 3: Object Tracking	<ul style="list-style-type: none"> Install <code>deep-sort-realtime</code> and integrate with the player tracker; ensure fallback IoU tracker still works. Calibrate Kalman filter parameters (process noise, measurement noise) for smooth ball trajectories with minimal lag. Tune optical flow parameters (window size, pyramid levels) for reliable motion estimation. Evaluate tracking quality: count ID switches per 100 frames, measure track continuity. Handle edge cases: ball occlusion by net, rapid direction changes, players crossing paths. Deliverables: tracking evaluation report (ID switches, continuity), demo clips, parameter docs, updated <code>object_tracking.py</code>. Knowledge: Kalman filter, Lucas-Kanade optical flow, DeepSORT, Hungarian algorithm, IoU matching. 	In progress

Member	Module	Responsibilities & Deliverables	Status
Nguyen Duc Dat	Module 4: 3D Trajectory & In/Out	<ul style="list-style-type: none"> After M1 homography fix: verify that ball court coordinates are plausible ($0 \leq x_c \leq 23.77$, $0 \leq y_c \leq 10.97$). Tune bounce detection: adjust velocity-change threshold, minimum inter-bounce time, confidence filtering. Validate in/out classification on 3–5 clips; report accuracy. Calibrate physics model parameters (initial height, drag coefficient) for realistic 3D trajectories. (Optional) Implement ML-enhanced bounce classification using trajectory features. Deliverables: bounce analysis report (per-clip results, confusion matrix), updated <code>trajectory_3d.py</code> and <code>in_out_classifier.py</code>. Knowledge: Projective geometry, physics modelling, Extended Kalman filter, classification. 	Blocked (depends on M1)
Bui Dinh Nguyen Minh	Module 5: Integration & Visualization	<ul style="list-style-type: none"> Integrate player foot-point projection onto mini-map (<code>player_court_positions</code>). Add court-line overlay on main video frame using H^{-1} projection of standard court key-points. Improve <code>CompositeFrameBuilder</code> layout: larger mini-map, score/stats panel. Package <code>run_pipeline.py</code> as a clean CLI entry point with progress bar and result summary. Produce final annotated demo video and output report. Write comprehensive README with installation steps, usage examples, and architecture diagram. Deliverables: final demo video, updated <code>pipeline.py</code> and <code>visualization.py</code>, README, run instructions. Knowledge: Software engineering, pipeline design, video I/O, data visualisation, documentation. 	In progress

9 Dependency Graph and Critical Path



Critical path: M1 (Homography fix) → M4 (3D + In/Out) → Final demo. The homography fix is the *highest priority* because Modules 4 and 5 cannot produce correct court-mapped results without a valid H .

Parallel work: M2 (Detection benchmarking) and M3 (Tracking stabilisation) can proceed independently of M1.

10 Milestones

1. **Milestone 1 – Homography Fix (Pham Hoang An):** Robust corner selection, reprojection sanity check, temporal stabilisation. Verified on multi-court test video.
2. **Milestone 2 – Detection Benchmark (Dao Duy Anh):** Comparative evaluation of all ball detection methods. Optimised default configuration. False-positive analysis.
3. **Milestone 3 – Stable Tracking (Duong Vu Duc):** DeepSORT integrated. ID switch count minimised. Ball trajectory smooth and continuous.
4. **Milestone 4 – Bounce/In-Out Validation (Nguyen Duc Dat):** Correct bounce detection on 3–5 validation clips. In/out accuracy report. Depends on Milestone 1.
5. **Milestone 5 – Final Demo (Bui Dinh Nguyen Minh + All):** Mini-map shows ball and players. Court overlay on video. Final annotated output video. Documentation and report submission.

11 Test Data and Validation

11.1 Primary Test Video

- File: data/raw/Video Project 4.mp4

- **Resolution:** 1920×1080 pixels
- **Frame rate:** 30 fps
- **Duration:** 1139 frames (≈ 38 seconds)
- **Scene:** Outdoor tennis facility, camera positioned above/behind one baseline. The main court is centre-left in the frame. At least one adjacent court is visible on the right, which causes interference in court detection.
- **Match type:** Doubles (4 players on the main court).

11.2 Latest Pipeline Results (Before Homography Fix)

- Total frames processed: 1139
- Processing time: ~ 636 s (1.8 fps on CPU)
- Ball detections: 633 / 1139 frames (55.6%)
- Player detections: 4556 total (≈ 4.0 per frame)
- Bounce events: 4 (all classified OUT — likely incorrect due to bad homography)
- Output file: `outputs/Video_Project_4_tracked.avi` (152 MB)

11.3 Validation Plan

- Collect 3–5 additional clips covering: single court (clean lines), multiple courts, partial visibility, different camera angles.
- Manually annotate ball positions and bounce events on selected frames for quantitative evaluation.
- Report: detection rate, tracking continuity, homography reprojection error, bounce classification accuracy.

12 Environment and Reproduction

12.1 Environment

- OS: Windows 10/11
- Python 3.13, PyTorch 2.6+
- Key packages: `ultralytics` (YOLOv8), `opencv-python`, `numpy`, `matplotlib`, `scipy`, `tqdm`, `pyyaml`
- Optional: `deep-sort-realtime` (for DeepSORT tracker)
- Environment variable: `TORCH_FORCE_NO_WEIGHTS_ONLY_LOAD=1`

12.2 Installation

```

1 pip install ultralytics opencv-python numpy matplotlib scipy tqdm pyyaml
2 pip install deep-sort-realtime # optional, for DeepSORT

```

12.3 Running the Pipeline

```

1 cd tennis-pickleball-tracker
2 set TORCH_FORCE_NO_WEIGHTS_ONLY_LOAD=1
3 python run_pipeline.py

```

12.4 Inspecting a Single Frame

```
1 python test_fix.py  
2 # Output: outputs/test_results/annotated_frame_v2.jpg
```

13 Conclusion

The Tennis/Pickleball Single-Camera Tracking System successfully implements an end-to-end pipeline from video input to annotated output. Ball detection achieves 55.6% frame-level detection rate, player filtering correctly identifies 4 on-court players per frame, and the tracking module produces continuous trajectories.

The primary remaining challenge is the **incorrect homography** caused by multi-court line interference (Section 6). This is the highest-priority fix because it blocks correct 3D reconstruction, bounce analysis, in/out classification, and mini-map rendering. The proposed fix (robust corner selection + reprojection sanity check + temporal stabilisation) is well-defined and assigned to Pham Hoang An.

With the 5-member task allocation defined in Section 8, the team is positioned to deliver a stable, evaluated system in the upcoming milestones.